



PUG TECH TALK · AUGUST 22, 2026

Building AI Agents with

Microsoft Foundry Local & Agent Framework

Manish Lad

AI Product Manager

HELLO — I'M

Manish Lad

AI Product Manager

- 12+ years of experience, 6+ working with AI
- 15+ years volunteering with Pune User Group
- I write about AI × product strategy — for builders and leaders.

[substack](#) manishlad.substack.com [in](#) [/manish-lad](#)

Why build **local-first**?

Why build local-first?

01

Data privacy & compliance

Process sensitive enterprise documents, healthcare data, or PII completely on-device. Zero network transmissions mean guaranteed compliance with air-gapped security requirements.

02

Zero latency & predictable cost

Eliminate cloud network round-trips and API rate limits. Achieve sub-10ms local token execution with standard zero recurring cloud API token costs.

Microsoft Foundry Local

Microsoft Foundry Local provides lightweight, high-performance local AI runtime for running open models (like Phi-4 and Mistral) directly on developer machines and edge hardware.

It exposes standard OpenAI-compatible endpoints seamlessly, enabling drop-in local replacements for cloud LLMs without altering agent code structure.

Microsoft Agent Framework

The Microsoft Agent Framework provides structured primitives for building reliable, task-oriented LLM workflows and multi-agent systems.

KEY CAPABILITIES

Deterministic state management and conversation routing.

Modular tool calling and structured output extraction.

Seamless binding to local inference runtimes.

Foundry Local runs the model. Agent Framework orchestrates **what you do with it.**

Agents, tool calling, conversation, workflows. Provider-agnostic — so local-first is a deployment decision, not an architectural commitment.

Demo

When to combine framework & local

01 Privacy-sensitive workflows

Triage, summarize, or extract structured metadata from confidential legal or medical documents locally.

02 Offline / edge deployment

Build intelligent agents that run on field equipment, vessels, or air-gapped industrial networks.

03 Cost reduction at scale

Replace repetitive bulk cloud LLM tasks with high-throughput local Small Language Models (SLMs).

04 Fast local prototyping

Iterate rapidly on agent tool definitions and prompt templates with zero latency and no cloud bill.

Q & A

Over to you.

THANK YOU

Let's connect.

Manish Lad

AI Product Manager

`substack` manishlad.substack.com

`linkedin` linkedin.com/in/manish-lad

Writing about AI × product strategy — for builders and leaders.



SCAN TO CONNECT

linktr.ee/manishlad